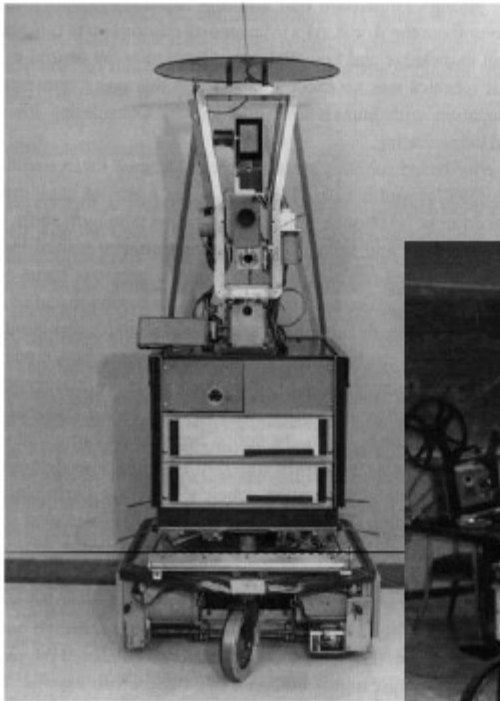
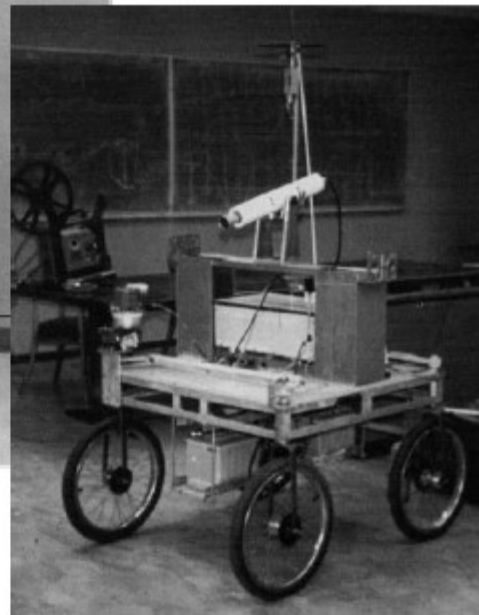


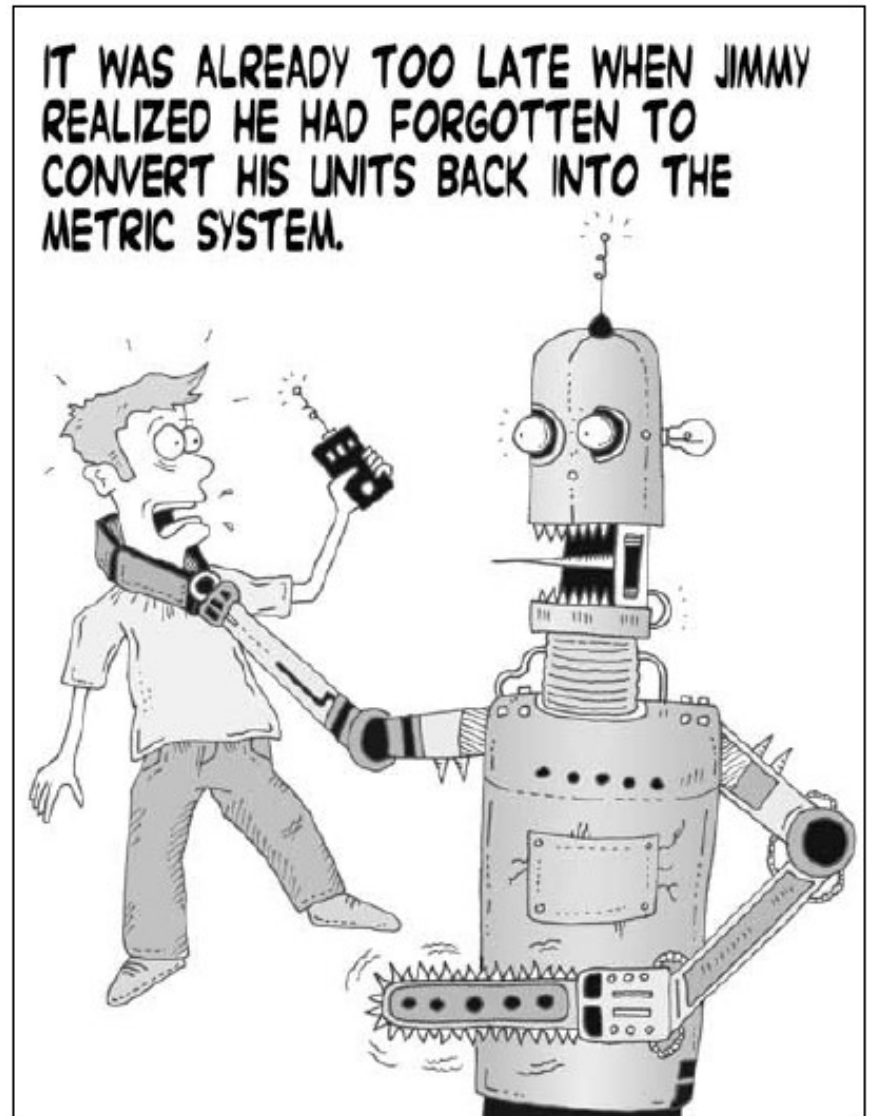
Hierarchical Paradigm



Shakey



Stanford Cart



IT WAS ALREADY TOO LATE WHEN JIMMY
REALIZED HE HAD FORGOTTEN TO
CONVERT HIS UNITS BACK INTO THE
METRIC SYSTEM.

Objectives

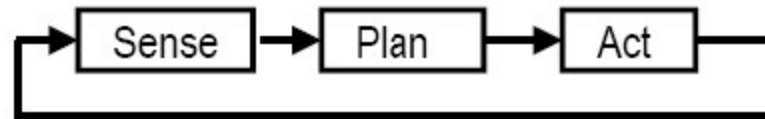
- To understand organizational differences between 3+ robot control paradigms
- To understand distinction between problem solver and planner
- To understand methods used by problem solvers
- To understand STRIPS approach to problem solving
- To understand example hierarchical architectures such as NHC and RCS

Robotic Paradigms

- A paradigm is a philosophy or set of assumptions and/or techniques which characterize an approach to a class of problems
- Applying the right paradigm makes problem solving easier
- There are generally **three paradigms** for organizing intelligence in robots: **hierarchical**, **reactive**, and **hybrid deliberative/reactive**

Three+ Primary Control Paradigms

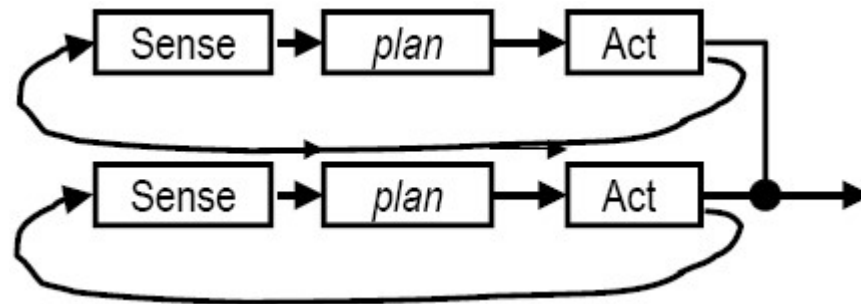
- Hierarchical



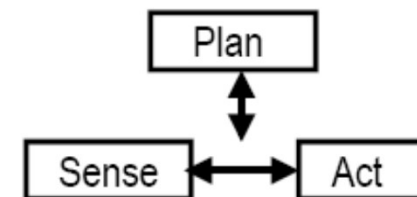
- Reactive



– Behavior-based

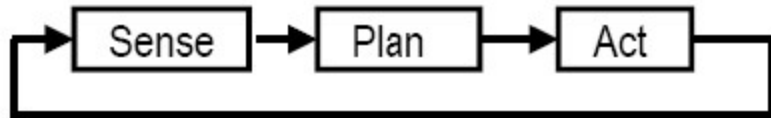


- Hybrid deliberative/reactive

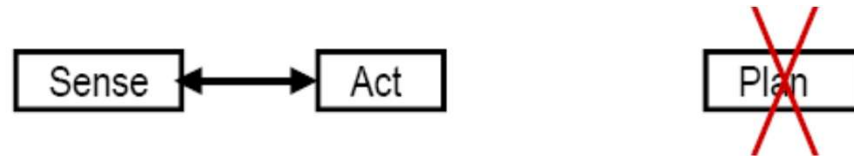


Hierarchical Paradigm

- 1967-1990
- Heavy on planning
- Senses the world, plans the next action, and acts
- Sensing data is gathered into one global world model
- World model is hard to maintain and brittle due to frame problem and closed world assumption



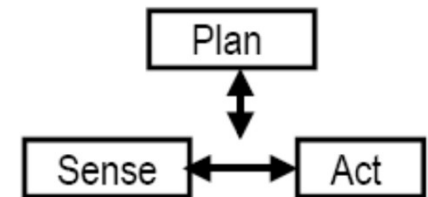
Reactive Paradigm



- 1988-1992
- Sensed information is directly coupled to an action, each is called a behavior
- Fast execution
- Motivated by:
 - Biology and cognitive psychology that examine living examples of intelligence
 - Low cost and increasing computing power on robot
- Multiple instances of SENSE-ACT couplings and combined results
- Throwing away planning is too extreme
- Behavior-based approach provides some levels of planning

Hybrid Paradigm

- 1990 - present
- Reactive paradigm serves as basis
- Robot first plans how to decompose a task into subtasks and then what behaviors to use for each subtask
- Then behaviors execute as per reactive paradigm
- Sensing is also hybrid, which connects to behavior directly and to the planner for constructing a task-oriented global world model



PLANNING IN HIERARCHICAL PARADIGM

Planning vs. Problem Solving

- Planning agent is very similar to problem solving agent
 - Constructs plans to achieve goals, then executes them
- Planning agent is different from problem solving agent in:
 - Representation of goals, states, actions
 - Use of explicit, logical representations
 - Way it searches for solutions

Problem Solver Characteristics

- Search-based problem solver:
 - **Representation of actions**: programs that generate successor state descriptions
 - **Representation of states**: complete state description; a data structure holding permutations of all possible states
 - **Representation of goals**: goal test and heuristic function to decide desirability
 - **Representation of plans**: solution is a sequence of actions; considers only unbroken sequences of actions

Famous Problem Solver Task: “Missionaries and Cannibals”

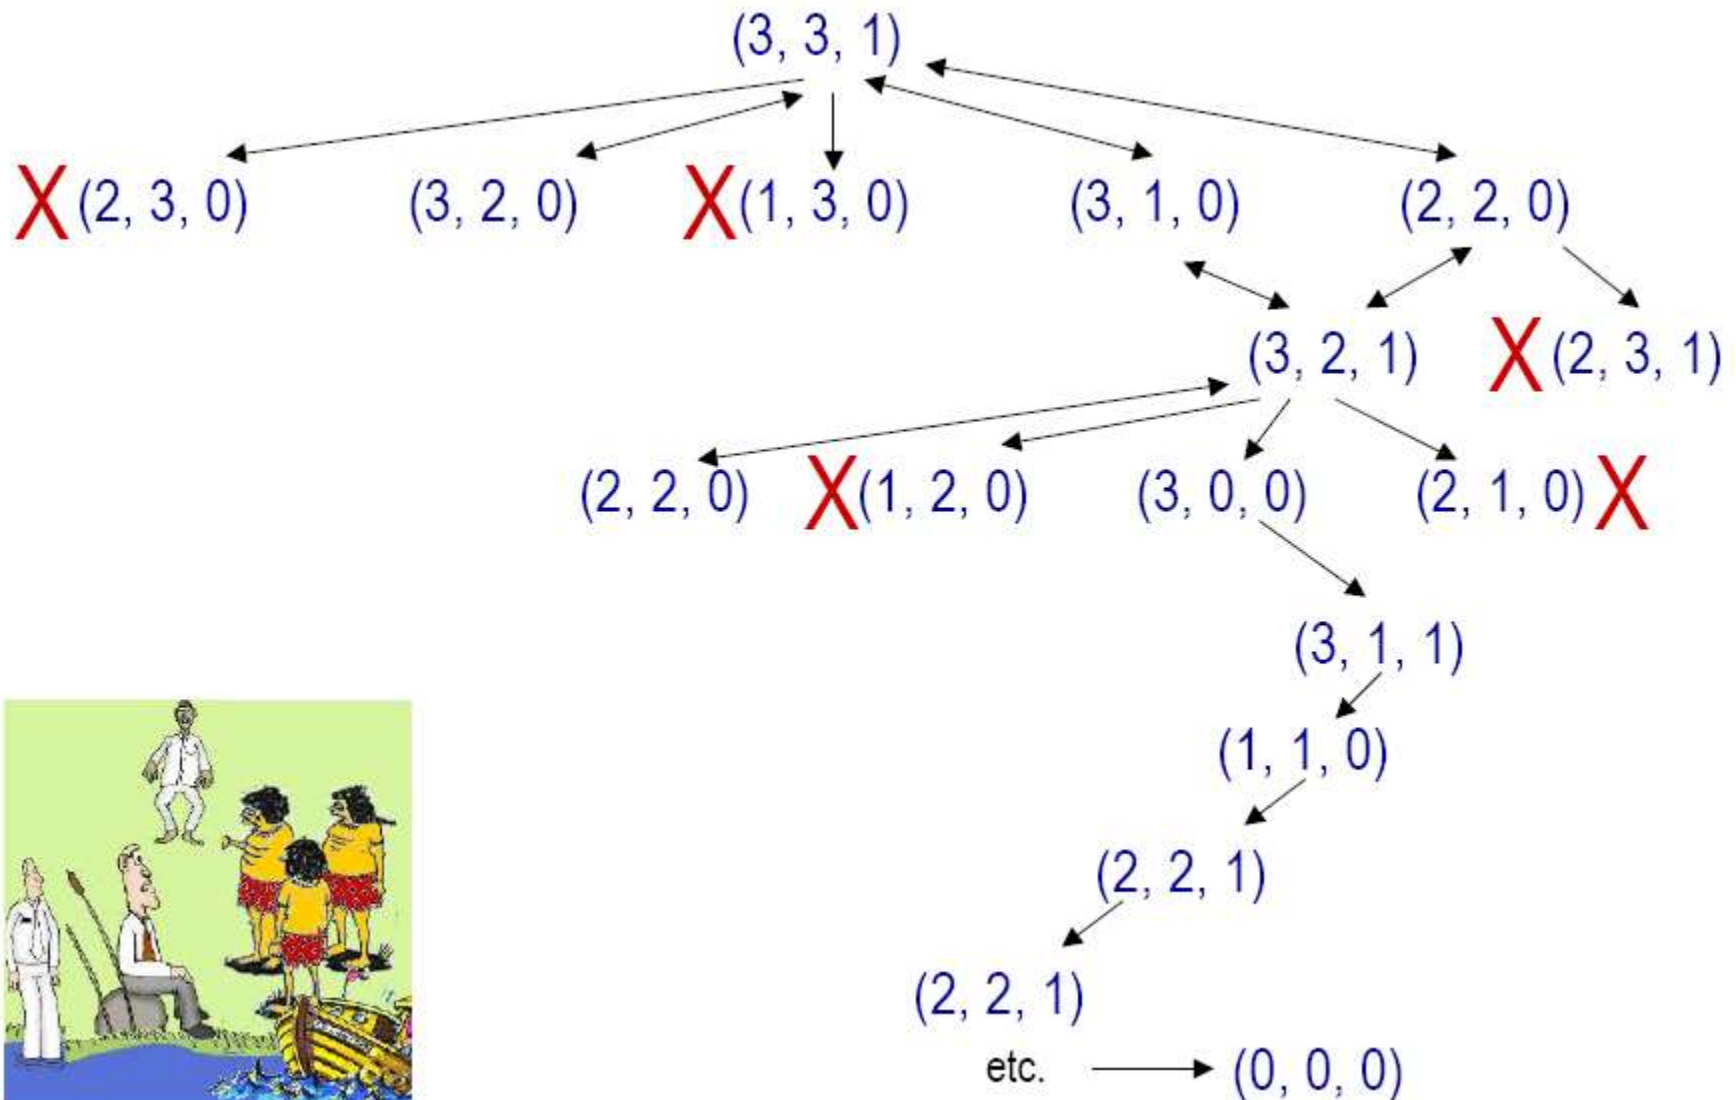
- “Missionaries and cannibals” problem:
 - Famous in AI
 - <http://www.learn4good.com/games/puzzle/boat.htm>



Formalizing Missionaries and Cannibals Problem

- Define states, actions (operators), goal test, path cost
 - **States:** $\langle m, c, b \rangle$ representing the # of missionaries and the # of cannibals, and the position of the boat
 - **Initial state:**
 - $\langle 3, 3, 1 \rangle$
 - **Actions:**
 - take 1 missionary, 1 cannibal, 2 missionaries, 2 cannibals, or 1 missionary and 1 cannibal across the river
 - **Goal test:**
 - $\langle 0, 0, 0 \rangle$
 - **Path cost:**
 - number of crossing

Solving Missionaries and Cannibals Problem

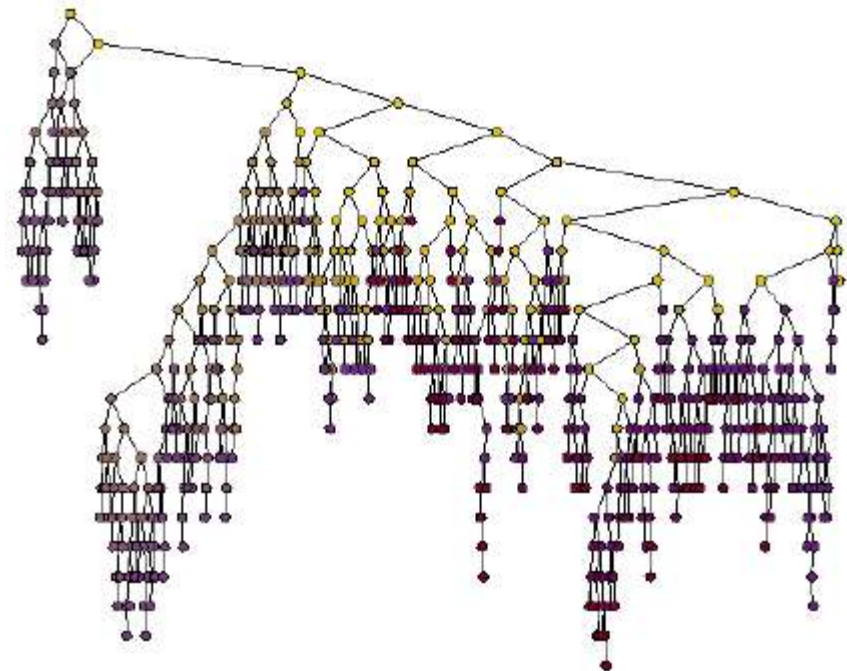


Search Strategies for Problem Solvers

- Key Criteria:
 - **Completeness**: guaranteed to find a solution when one exists
 - **Time complexity**: how long does it take to find solution?
 - **Space complexity**: how much memory is needed to find solution?
 - **Optimality**: does strategy find the highest-quality solution?

Various Search Strategies for Problem Solver

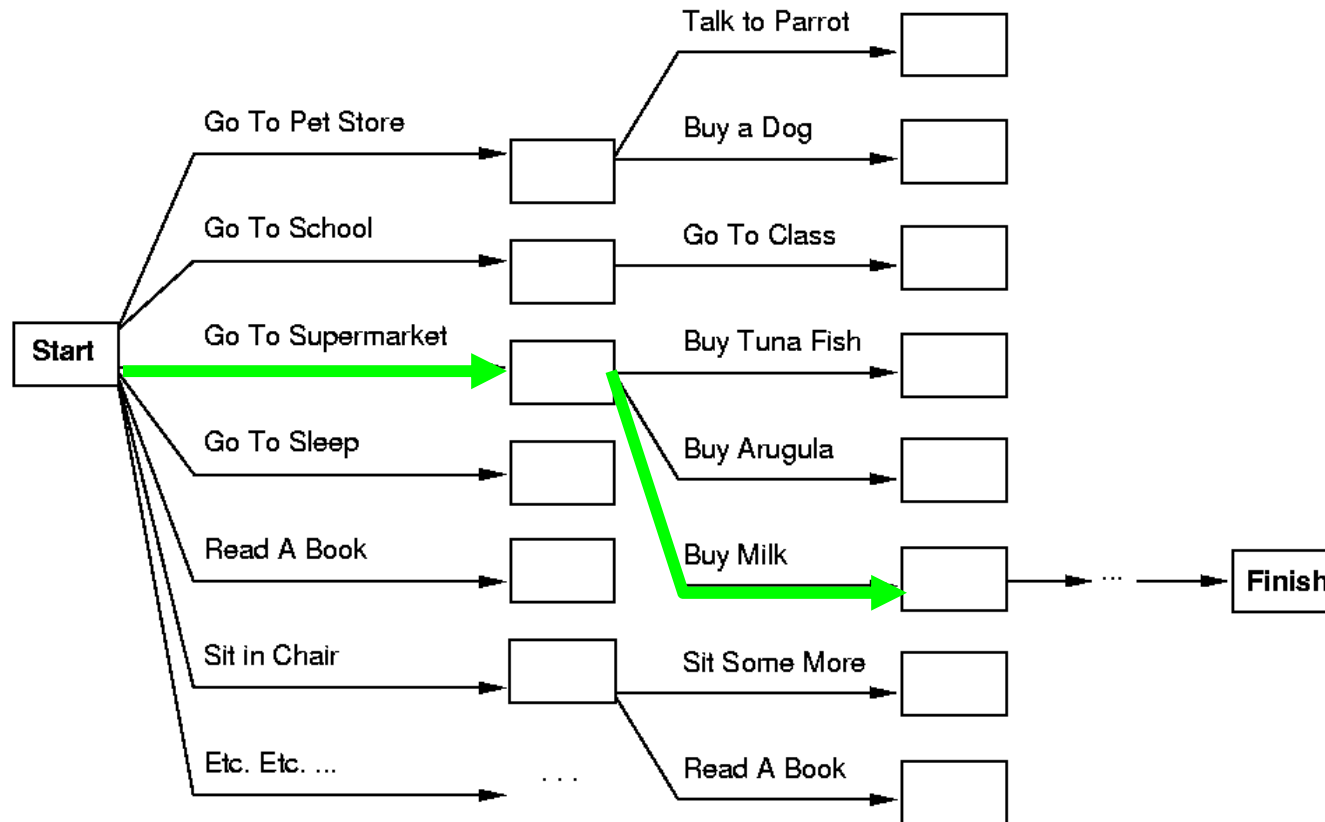
- **Uninformed search** (no info. about path cost from current state to goal):
 - Breadth-first
 - Uniform cost
 - Depth-first
 - Depth-limited
 - Iterative deepening
 - Bidirectional
 - Etc.
- **Informed search**
 - Greedy
 - A*
 - Hill-climbing/gradient descent
 - Simulated annealing
 - Etc.



(a messy search tree)

Search vs. Planning

- Consider the task: get milk, bananas, and a cordless drill
 - Standard search algorithms seem to fail miserably
 - Why? Huge branching factor & heuristics



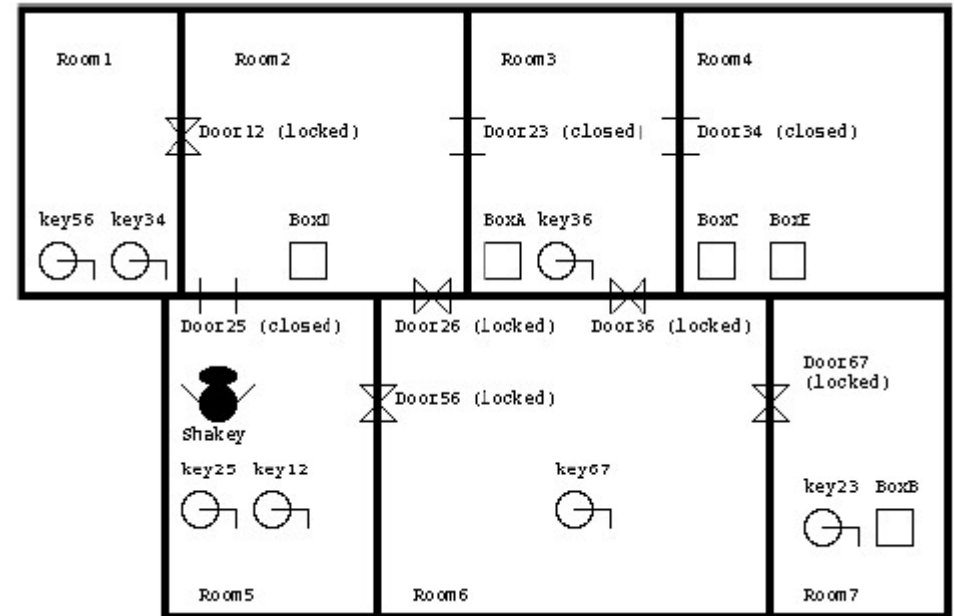
Search vs. Planning

- Planning systems do the following:
 - divide-and-conquer
 - relax requirement for sequential construction of solutions

	Search	Planning
States	data structures	logical sentences
Actions	code	preconditions/outcomes
Goal	code	logical sentences
Plan	sequence from s_0	constraints on actions

Planning-Based Approach to Robot Control

- **Job of planner:** generate a goal to achieve, and then construct a plan to achieve it from the current state
- Must define representations of:
 - **Actions:** generate successor state descriptions by defining preconditions and effects
 - **States:** data structure describing current situation
 - **Goals:** what is to be achieved
 - **Plans:** solution is a sequence of actions

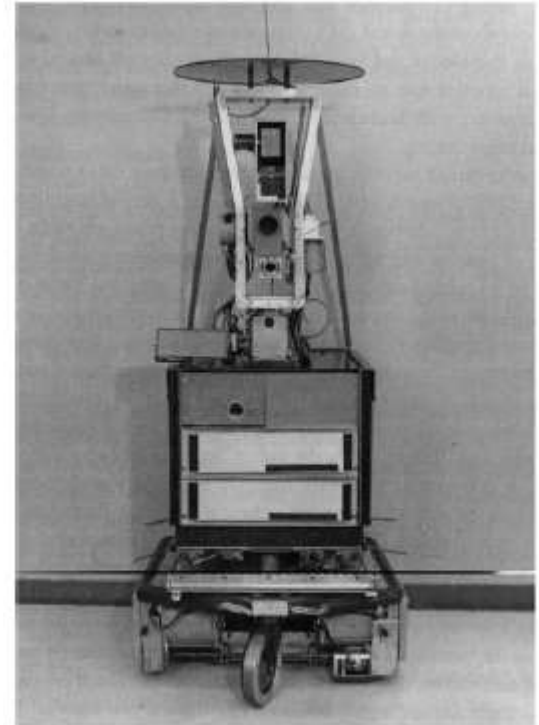


First Order Predicate Calculus

- First order predicate calculus: formal language useful for making inferences and deductions
- Elementary components:
 - Predicate symbols (e.g., WRITES(), OWNS(), LIVES())
 - Variable symbols (e.g., x , y)
 - Function symbols (e.g., father(x))
 - Constant symbols (e.g., HOUSE-1)
 - Connectives: and, or, negation, implies
 - Quantification:
 - Universal
 - Existential

STRIPS-Based Approach to Robot Control

- Use first-order logic and theorem proving to plan strategies from start to goal
- STRIPS language:
 - “Classical” approach that most planners use
 - Lends itself to efficient planning algorithms
- Cost: \$100, 000
- See a video of Shakey here:
<https://www.youtube.com/watch?v=GmU7SimFkpU&t=46s>



Shakey (SRI), 1960's

STRIPS

- STRIPS (STanford Research Institute Problem Solver)
 - a restrictive way to express states, actions and goals, but leads to more efficiency
- **States**: conjunctions of ground, function-free, and positive literals
 - E.g. $\text{At}(\text{Home}) \wedge \text{Have}(\text{Banana})$
- **Goals**: conjunctions of literals, may contain variables (existential), hence goal may represent more than one state
 - E.g. $\text{At}(\text{Home}) \wedge \neg \text{Have}(\text{Bananas})$
 - E.g. $\text{At}(x) \wedge \text{Sells}(x, \text{Bananas})$
- **Actions**: preconditions that must hold before execution and the effects after execution

STRIPS Action Schema

- An **action schema** includes:
 - **Action name & parameter list** (variables): name for what an agent does
 - **Precondition**: a conjunction of function-free positive literals. Any variables in it must also appear in parameter list
 - **Effect**: a conjunction of function-free literals (positive or negative)
 - ADD List
 - DELETE List

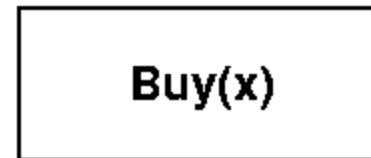
- Example:

Action: Buy (x)

Precondition: At (p), Sells (p, x)

Effect: Have(x)

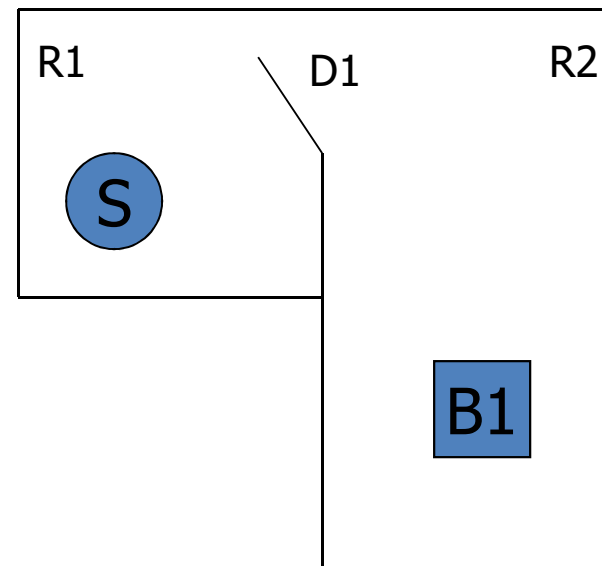
At(p) Sells(p,x)



Have(x)

An Example of STRIPS World Model

- A world model is a set of facts
- The robot's knowledge can be represented by the following predicates:
 - **INROOM**(*x*, *r*), where *x* is a movable object, *r* is a room
 - **NEXTTO**(*x*, *t*)
 - **STATUS**(*d*, *s*), where *d* is a door, *s* means OPEN or CLOSED
 - **CONNECTS**(*d*, *rx*, *ry*)
- The world model in this figure can be represented with the above predicates, with the initial state and goal state



The Actions in STRIPS World

- Types of actions Shakey can make
 - Move from place to place:
 - **GOTODOOR(S, dx):**
 - PRECOND: $\text{INROOM}(S, rk) \wedge \text{CONNECT}(dx, rk, rm)$
 - Effect:
 - » add-list: $\text{NEXTTO}(S, dx)$
 - » delete-list: null
 - **GOTHRUDOOR(S, dx):**
 - PRECOND: $\text{CONNECT}(dx, rk, rm) \wedge \text{NEXTTO}(S, dx) \wedge \text{STATUS}(dx, \text{OPEN}) \wedge \text{INROOM}(S, rk)$
 - Effect:
 - » add-list: $\text{INROOM}(S, rm)$
 - » delete-list: $\text{INROOM}(S, rk)$

How to Develop a Plan in STRIPS?

- Use “means-ends analysis”
 - If robot is not at goal, then **measure the distance** (or difference) between robot’s current state and its goal state
 - Find an **action** that the robot can take that will reduce this distance (difference), and select it if open conditions are satisfied
 - Make the first false precondition of this action the new “**subgoal**”, remembering old goal by pushing on stack
 - Recursively reduce difference by repeating above
 - When all open conditions for an action are satisfied, **add the action to the plan stack** and update world model
 - When done, pop the actions off the stack to create the **plan** from start to goal

Go Back to the Previous Example

For details, please read the textbook
2.2.2.

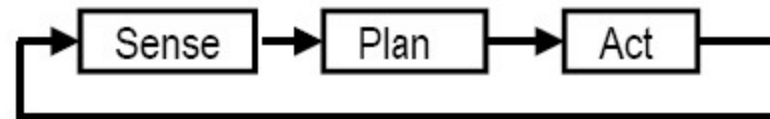
Summary of Planning Problem

- **Planning agents** use look-ahead to find actions to contribute to goal achievement
- Planning agents differ from problem solvers in their use of **more flexible representation** of states, actions, goals, and plans
- The **STRIPS language** describes actions in terms of preconditions and effects

Hierarchical Robotic Control: Representative Architectures

Recall: Hierarchical Paradigm

- Hierarchical paradigm



- The hierarchical paradigm is sequential and orderly
 - First the robot senses the world and constructs a global **world model**
 - Then “eyes” closed, the robot **plans** all the directives needed to reach the goal
 - Finally, the robot **acts** to carry out the first directive
 - After the robot has carried out the whole sequence, it repeats the cycle again

World Model

- All sensor observations are fused into one global data structure – the **world model**
 - a priori representation of the environment the robot is operating in (e.g., map)
 - sensing information (e.g., “I am at a doorway”)
 - any additional cognitive knowledge
- Recall our STRIPS planner
 - it works recursively by backward chaining
 - it generates a list of operators (actions) to apply
 - it plans first and then executes

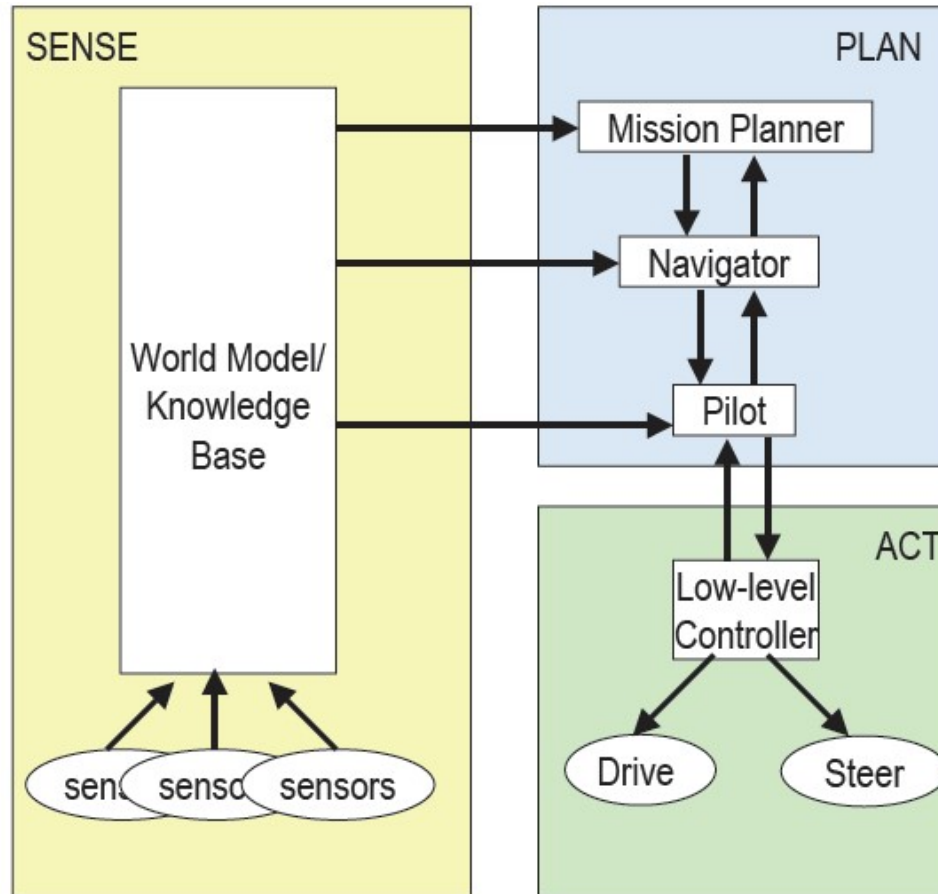
Architectures

- “Architecture”: a method of implementing a control paradigm, of embodying certain principles in a concrete way
- Evaluation of existing architecture is based on the following criteria:
 - Support for modularity
 - does it show good software engineering principles?
 - Niche targetability
 - how well does it work for intended applications?
 - Ease of portability to other domains
 - how well would it work for other applications or other robots?
 - Robustness
 - is the system vulnerable?

Case Studies

- Two best-known hierarchical architectures:
 - Nested Hierarchical Controller (NHC)
 - developed by Meystel
 - NIST Realtime Control System (RCS)
 - adapted to a teleoperation version called NASREM
 - developed by Albus

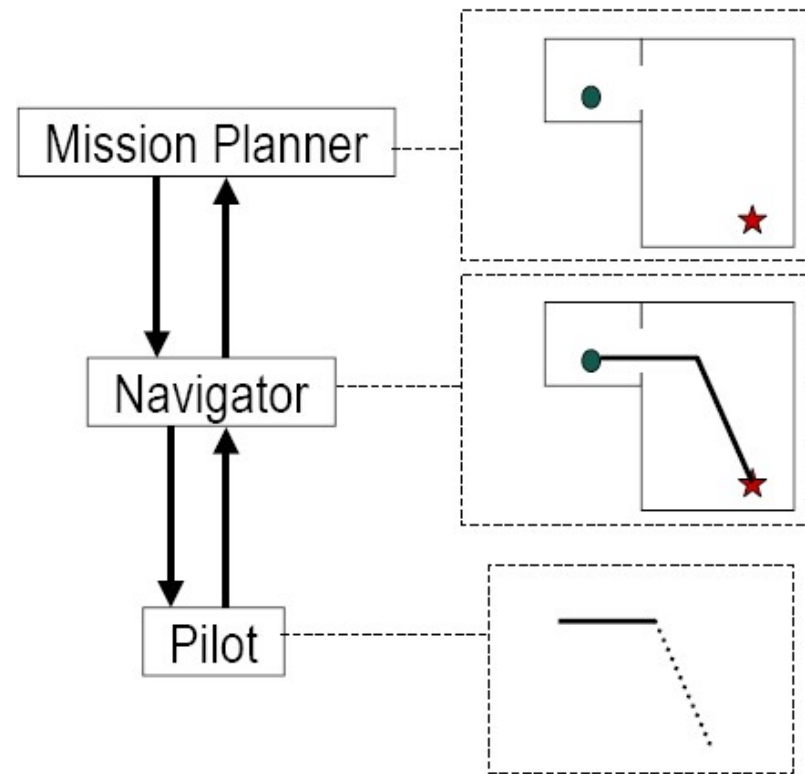
Nested Hierarchical Controller



Major contribution of NHC: decomposition of planning into three subsystems

Planning is Hierarchical

- Uses map to locate self and goal
- Generates path from current position to goal
- Generates actions robot must execute to follow path segment



Details of NHC

- **Mission planner:**
 - either receives a **mission** from a human or generates a mission for itself
 - responsible for **translating** the mission into terms that other functions can understand (e.g., access a map of the building and locates where the robot is and where the goal is)
- **Navigator:**
 - takes information from mission planner (e.g., current position, goal position) and **generates a path** in the form of a set of waypoints, or straight lines for the robot to follow
- **Pilot:**
 - takes the first straight line or path segment and determines what **actions** the robot has to do to follow the path segment (e.g., speed, turning rate)

Details of NHC, Continue

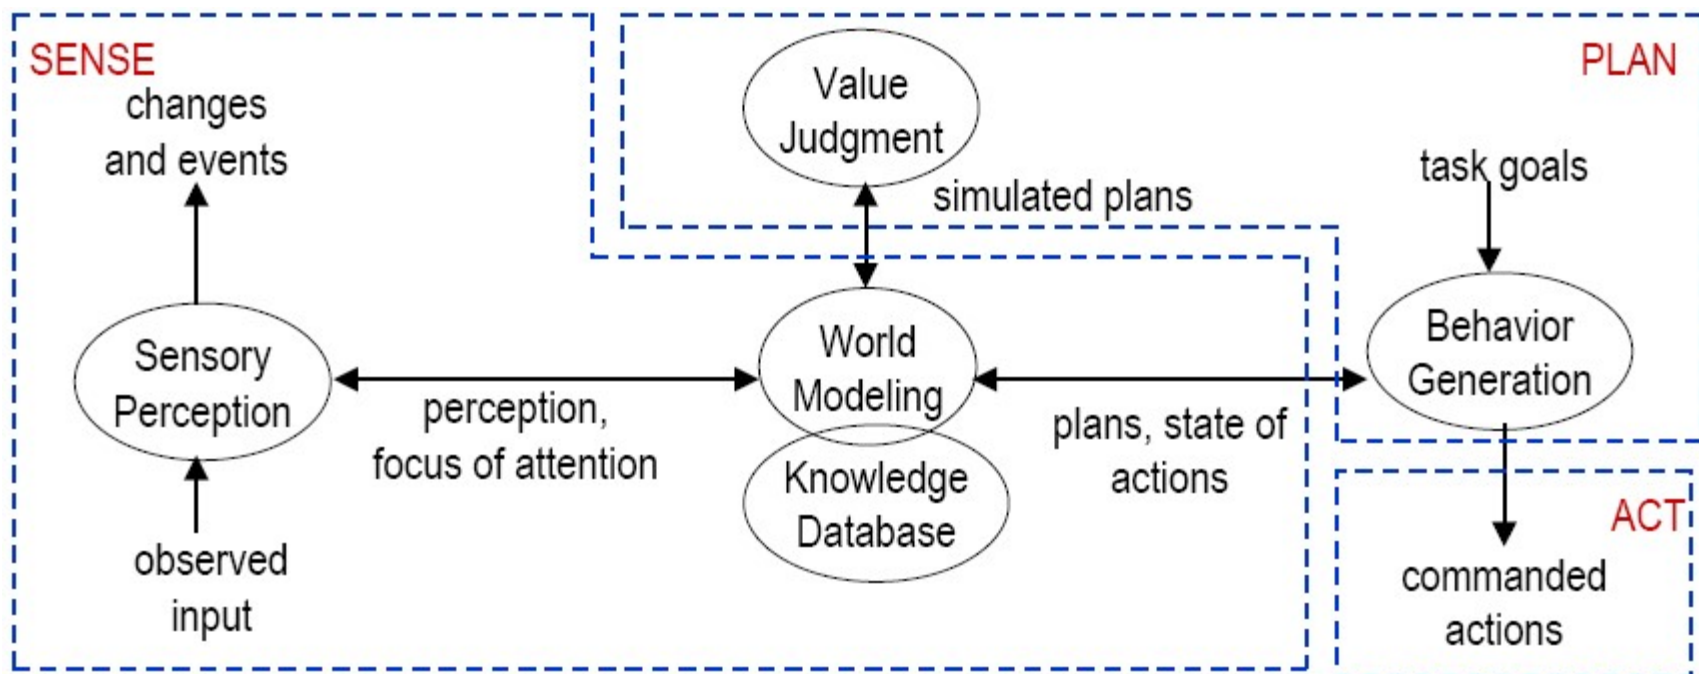
- After the robot has executed the commands from Pilot, it senses again and updates the world model
 - the entire planning cycle does not need to be repeated
 - it checks the changes, and repeats the corresponding part
 - e.g., if the robot has reached its waypoint, the Pilot informs the Navigator; if the waypoint isn't the goal, then pass the next waypoint for the robot to reach; if the waypoint is the goal, then the Navigator informs the Mission Planner; the Mission Planner may issue a new goal, etc.
 - e.g., if the robot has encountered an obstacle to its path, the Pilot passes control back to the Navigator; the Navigator computes a new path, and gives it to the Pilot to carry out

Advantage/Disadvantage of NHC

- **Advantage:**
 - Interleaves planning and action (unlike STRIPS)
 - Plan is changed if world is different from expected
 - Note that the decomposition is inherently hierarchical in intelligence and scope
 - Mission planner is smarter than Navigator, who is smarter than the Pilot
- **Disadvantage:**
 - Planning decomposition is only appropriate for navigation tasks

NIST RCS & NASREM

- NIST RCS (National Institute of Standards and Technology **Real-time Control System**): developed to serve as standard for manufactures' development of more intelligent robots
- Similar in design to NHC
 - **Primary difference**: sensory perception includes preprocessing (feature extraction, or focus of attention)



Details of RCS

- **Value judgment** plans and the simulates the plans to ensure they will work
- The planner gives the plan to **Behavior Generation** module, which converts the plans into actions

Multiple Layers of NIST RCS & NASREM

- Each layer has:
 - Sensory processing
 - World modeling
 - Task decomposition
 - Value judgment
- All layers joined by global memory through which representational knowledge is shared
- Perception is not tied directly to action
- Perception is still integrated into a global world model

NHC and RCS/NASREM: Well-Suited for Semi-Autonomous Control

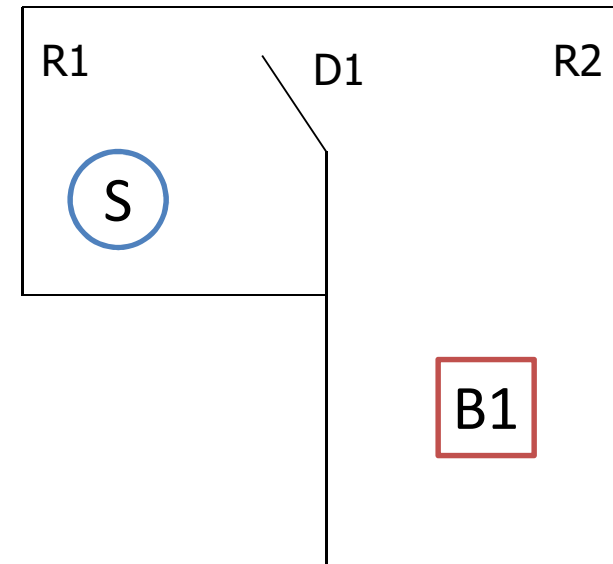
- Human operator could:
 - Provide world model
 - Decide mission
 - Decompose mission into plan
 - Decompose plan into actions
- Lower-level controller (robot) could:
 - Execute actions

Evaluations of Hierarchical Architectures

- Robots (other than the Walter's tortoise and Braitenberg's vehicles) built before 1985 typically used hierarchical style
- Primary advantage:
 - provides an ordering of the relationship between [sensing, planning, and acting](#)
- Primary disadvantage:
 - planning: every update cycle, robot had to [update](#) global world model, then plan, which introduces a bottleneck
 - sensing and acting are [disconnected](#)
 - appropriate hierarchical decomposition is [application-dependent](#)
 - [uncertainty](#) not well handled

Exercise: STRIPS Example

- Predicates:
 - **INROOM**(x, r), where x is a movable object, r is a room
 - **NEXTTO**(x, t), where x is a movable object, t is a door or movable object
 - **STATUS**(d, s), where d is a door, s means OPEN or CLOSED
 - **CONNECTS**(d, rx, ry)



- Difference table

Operator	Preconditions	add-list	delete-list
OP1: GOTODOOR(S, dx)	INROOM(S, rk) CONNECT(dx, rk, rm)	NEXTTO(S, dx)	
OP2: GOTHRUDOOR(S, dx)	CONNECT(dx, rk, rm) NEXTTO(S, dx) STATUS($dx, OPEN$) INROOM(S, rk)	INROOM(S, rm)	INROOM(S, rk)

Exercise

- If I want the robot to move from the position that's next to door to box B1 in the same room
 - What operator/predicates do we need to add?
 - Construct a plan
 - Show the changes in the world model after each operator is applied

Partial Answer

GotoBox(S, bx)

Preconditions:

INROOM(S, rk), INROOM(bx, rk),

NEXTTO(S, dx)

Add-list:

NEXTTO(S, bx)

Delete-list:

NEXTTO(S, dx)